



Kafka Monitoring: What Matters!

Amrit Sarkar

THIS IS NOT A CONTRIBUTION

Agenda

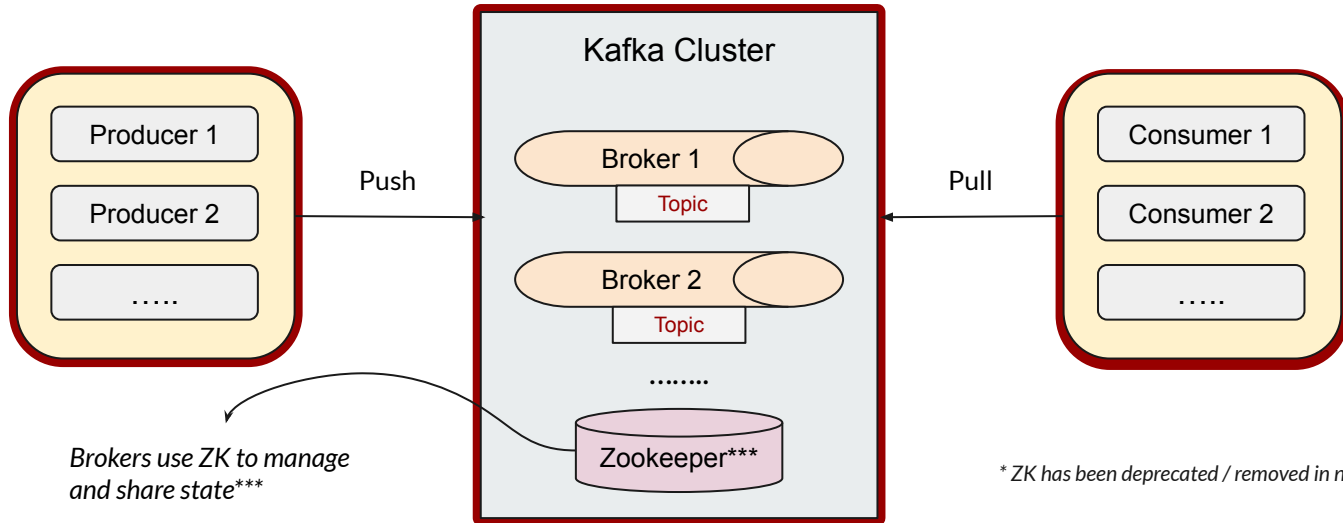


- Kafka Basics
- Performance Areas
- Need for Observability
 - Monitoring Options
- Performance classification around Components
- Kafka Consumer Lag evaluation
 - absolute to relative
- Trend Analysis

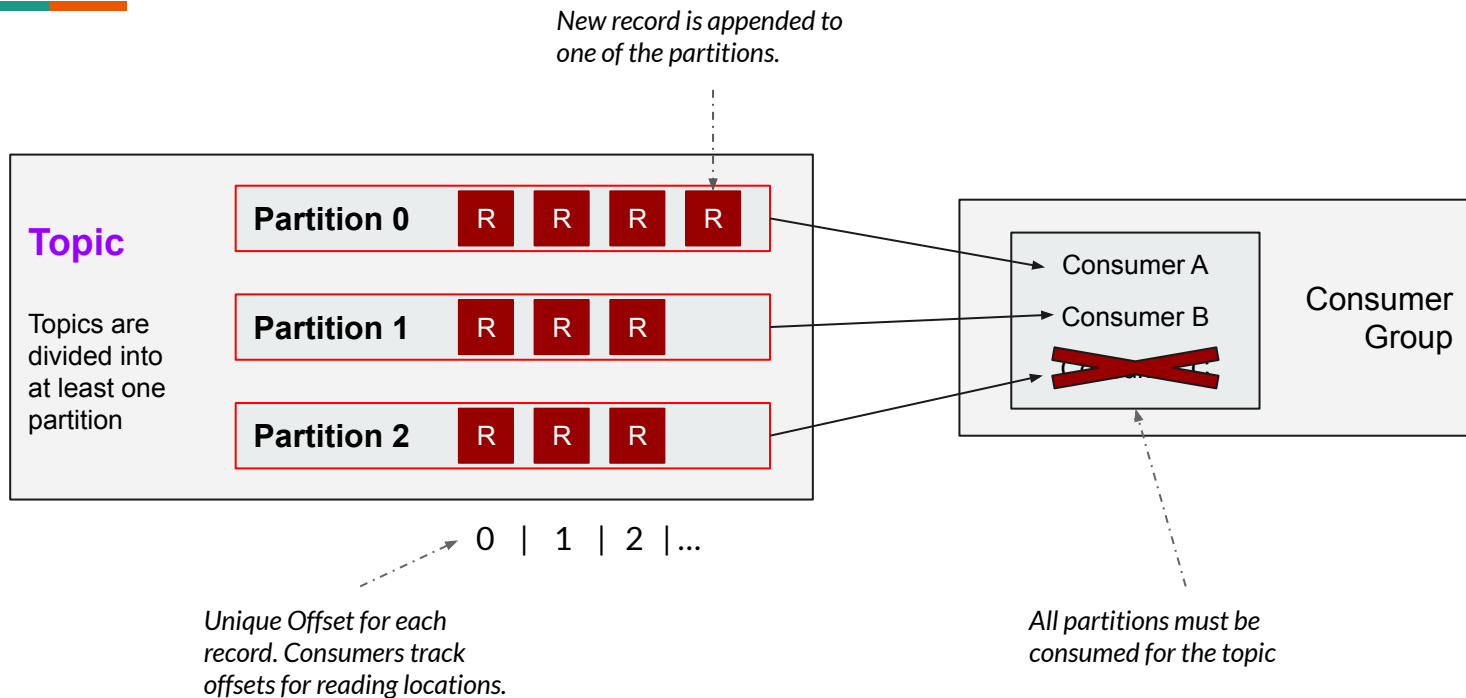
Kafka Basics

Kafka moves data between producers (writers) and consumers (readers),
with data protection, high availability, low latency at high scale!

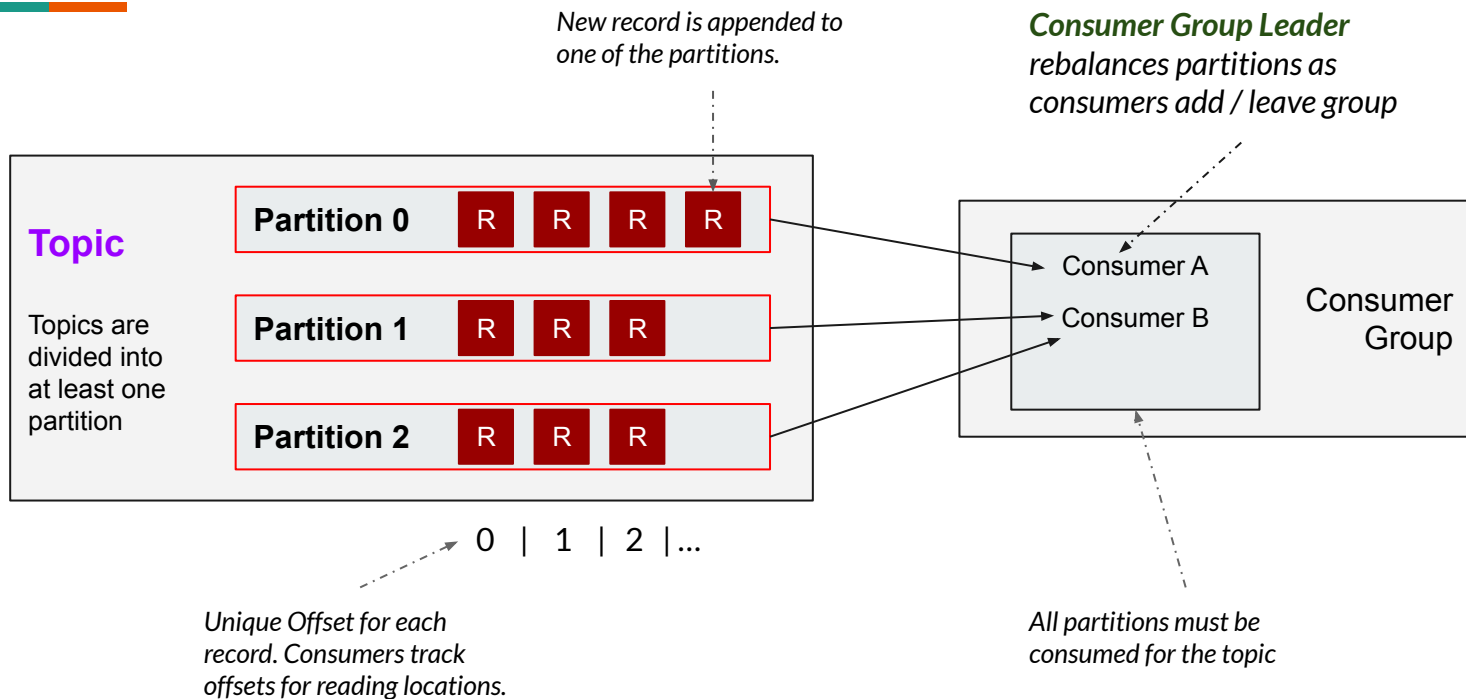
Use cases: Metrics, Log Aggregation Solution & Stream Processing



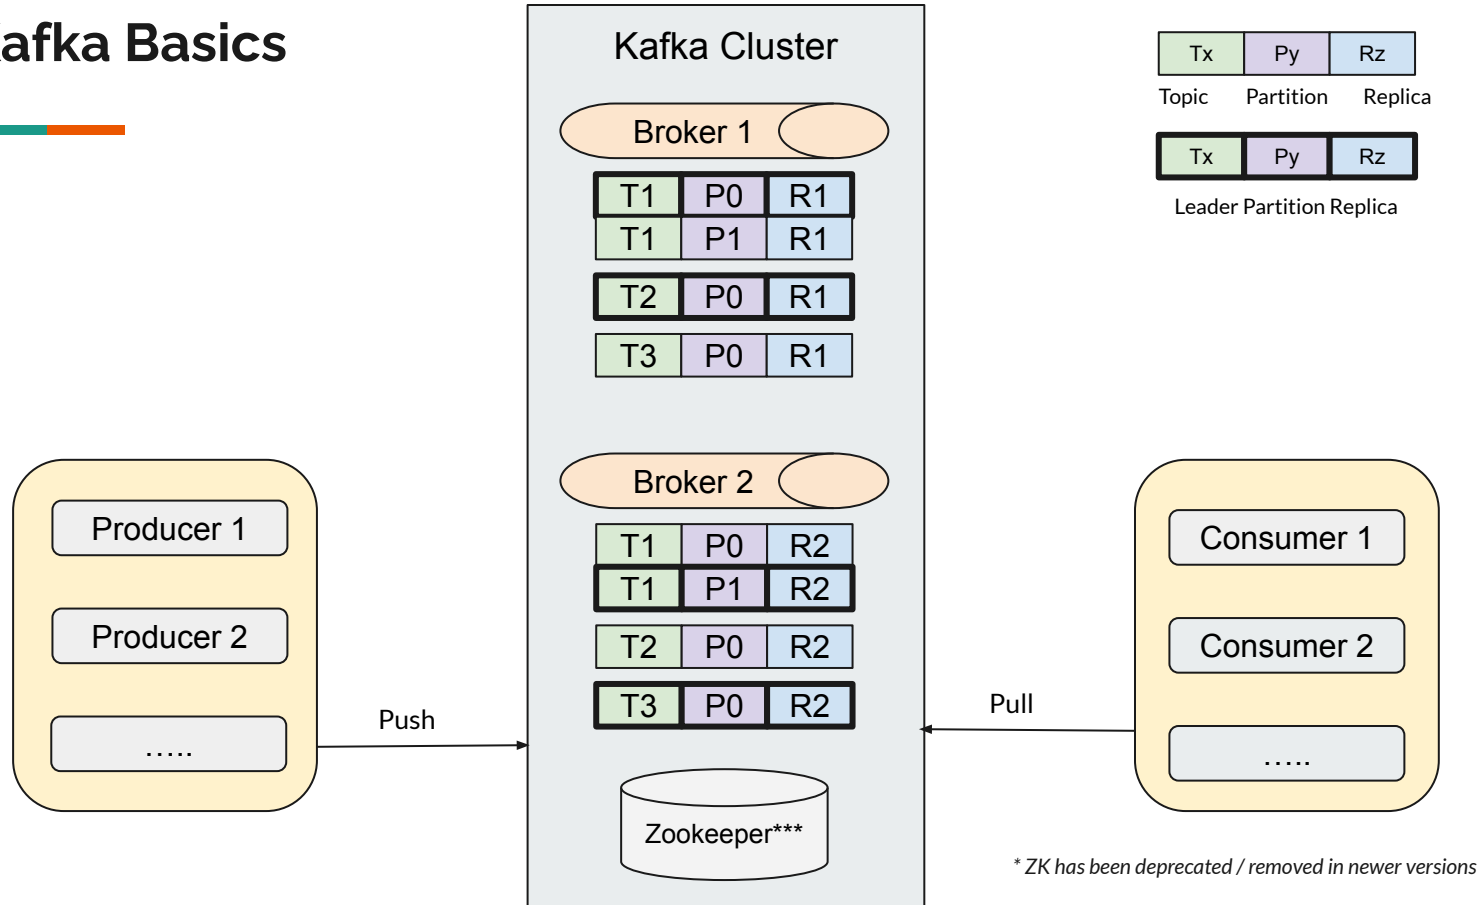
Kafka Basics



Kafka Basics



Kafka Basics

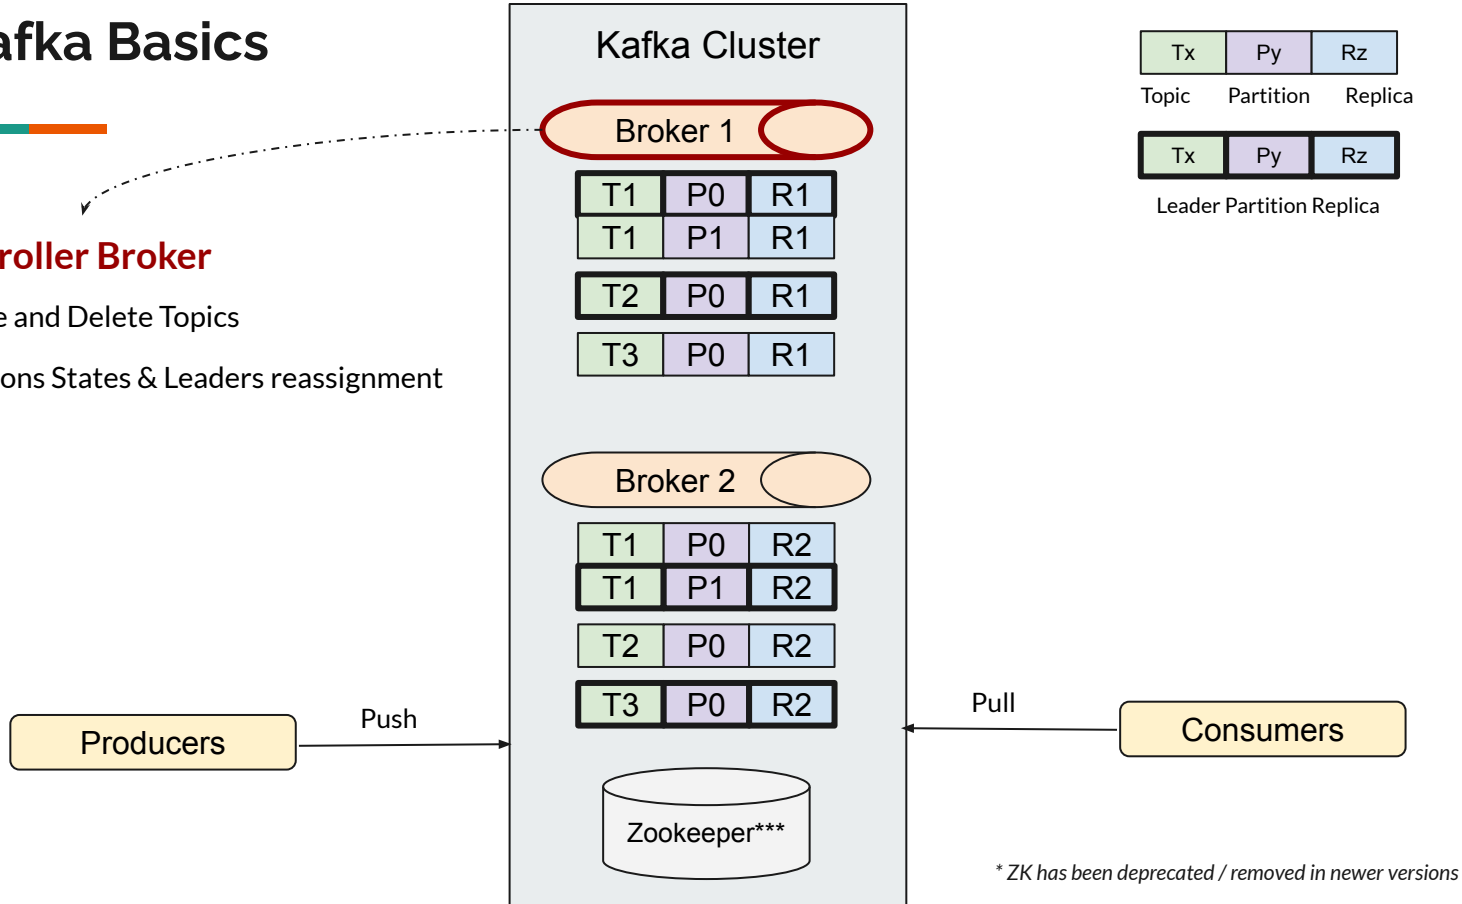


Kafka Basics



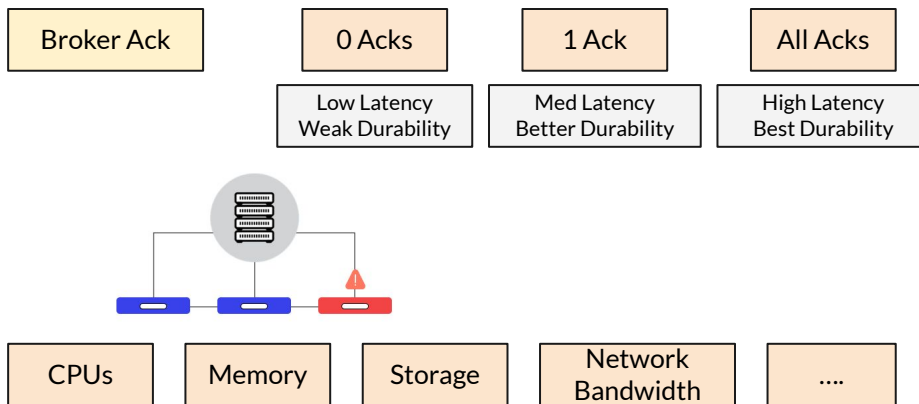

Controller Broker

- * Create and Delete Topics
- * Partitions States & Leaders reassignment



Performance Areas

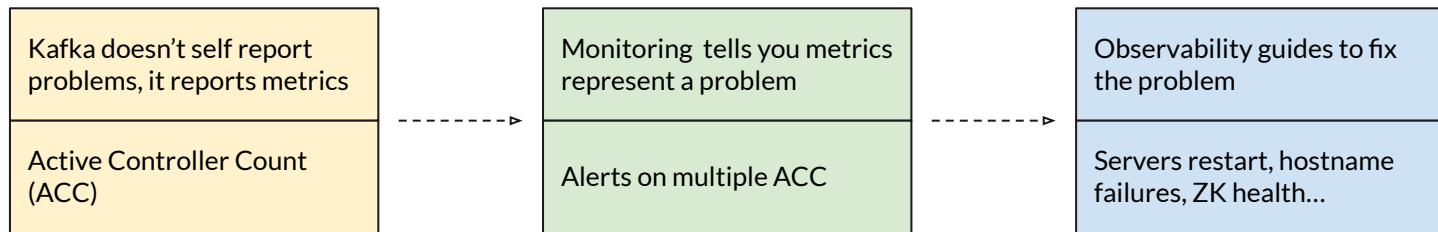
- Throughput & Latency
 - Production Rates
 - Consumption Rates
 - Consumers' Lag
- Data Integrity
 - Reads Confirmation
 - Writes Confirmation
- Fault Tolerance
 - No business impact on failure
- Resource Usage



Why do we need Observability?

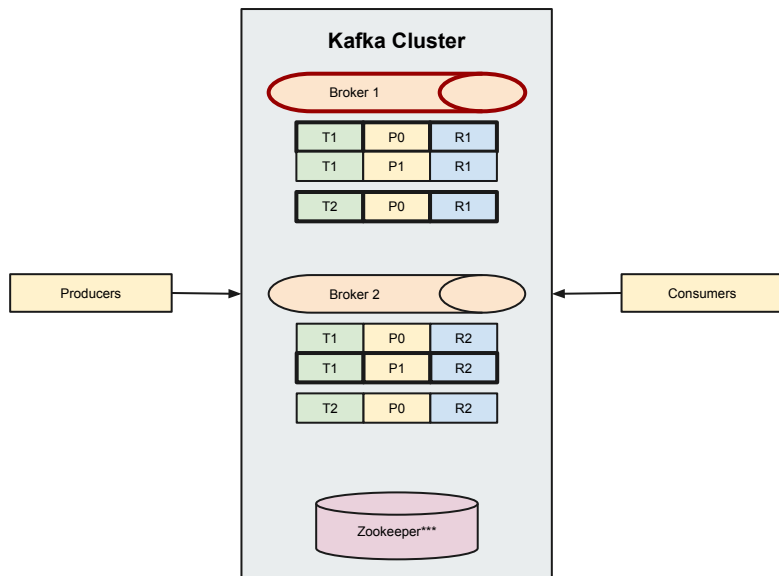
Pre-built dashboards monitor and alert for anticipated future performance issues.

Explore and quickly identify unanticipated issue root causes in an observability scenario.



Performance classification around Components

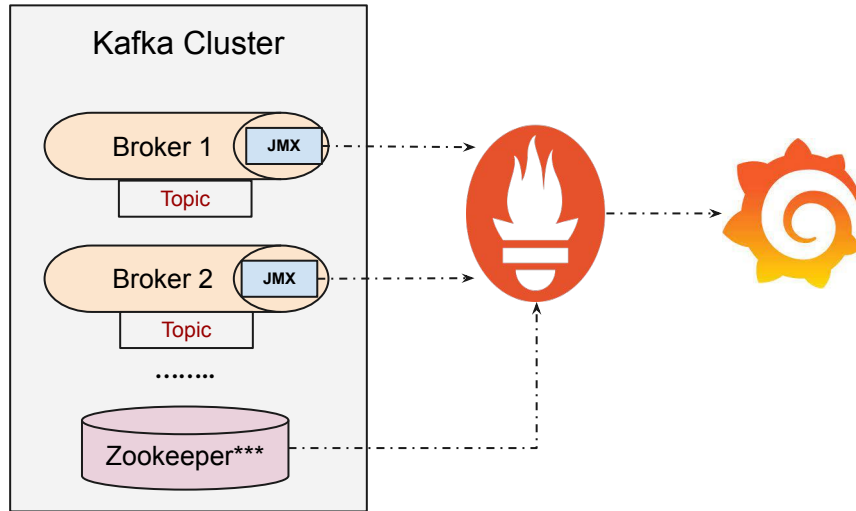
Each component act as potential factor in the performance of Kafka messaging system.



* ZK has been deprecated / removed in newer versions

- Producers
 - Rate
- Brokers
 - Topics Health
 - Load Distribution
- Topics
 - Partition Health
 - Load Distribution
- Consumers
 - Rate
 - **Lag**
- Generic
 - Transmission (Network) health
 - Capacity

Monitoring Options - Getting Metrics In



- Confluent Control Center
- KafDrop
- Yahoo Kafka Manager
- Cruise Control
- Kafka Monitor
- Kafka Tool

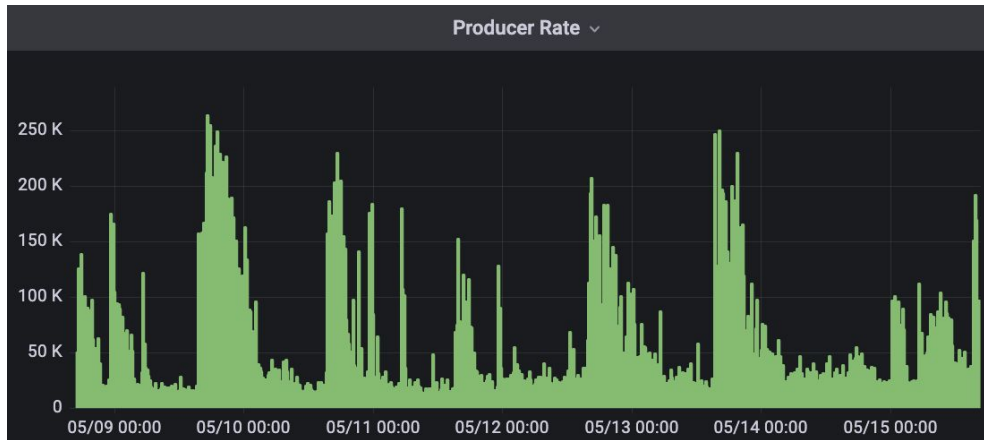
... and more

Producer: Rate



- **send()** function call (or similar) to push data → [RecordMetadata](#) object
 - **offset()** function returns a **LONG** → offset of the the record in the topic-partition.

Offset value can be pushed to metrics store for visualisation.



Producer: Compression & Latency



- Bigger batches →
 - higher throughput
 - less compression
 - Small enough to keep GC ⚡
($< 10\text{mb}$)

Ideally

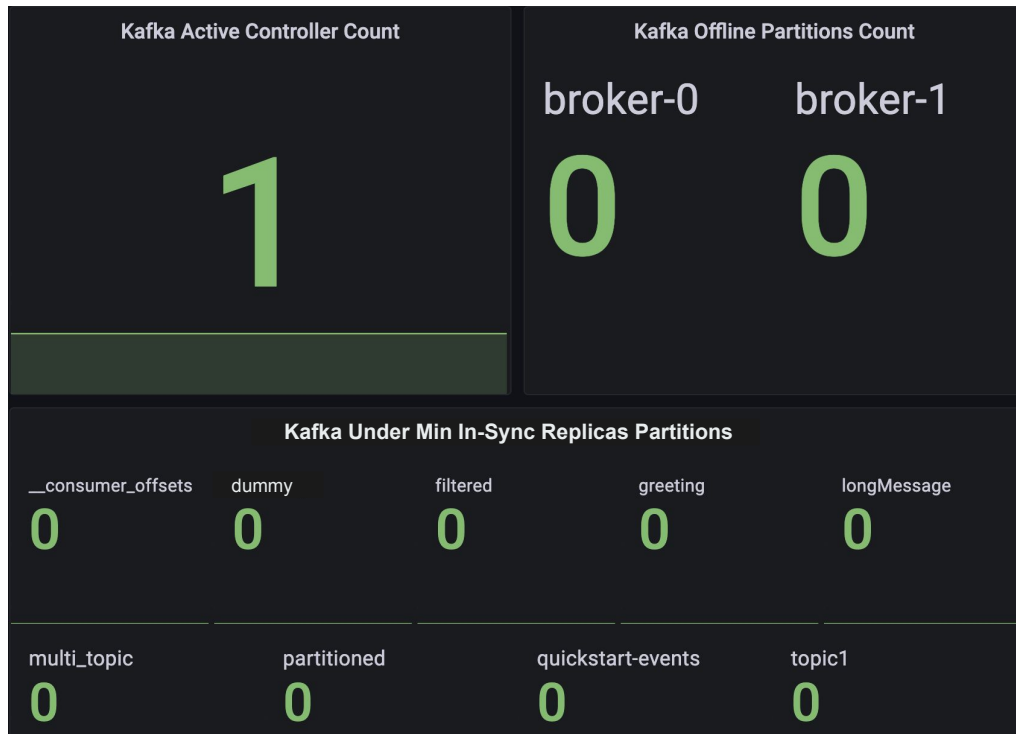
- Batch Size in Bytes → **Optimally High**
*kafka.producer:type=producer-metrics,client-id="{client-id}"
batch-size-avg*
- Compression Rate → **LOW**
*kafka.producer:type=producer-metrics,client-id="{client-id}"
compression-rate-avg*
- Request Latency → **LOW**
*kafka.producer:type=producer-metrics,client-id="{client-id}"
request-latency-avg*

The Big Four - Key Metrics (JMX)



- Number of active controllers, **must be = 1**
*kafka.controller:name=ActiveControllerCount,
type=KafkaController*
- Number of under min ISR partitions, **must be = 0**
*kafka.server:name=UnderMinIsrPartitionCount,
type=ReplicaManager*
Checkout: 'UnderReplicatedPartitions' metric too
- Number of offline partitions, **must be = 0**
*kafka.controller:name=OfflinePartitionsCount,
type=KafkaController*
- Consumer Lag (per partition)
*kafka.consumer:name=MaxLag,
type=ConsumerFetcherManager,clientId={[-.w]+}*

The Big Four - Key Metrics (JMX)



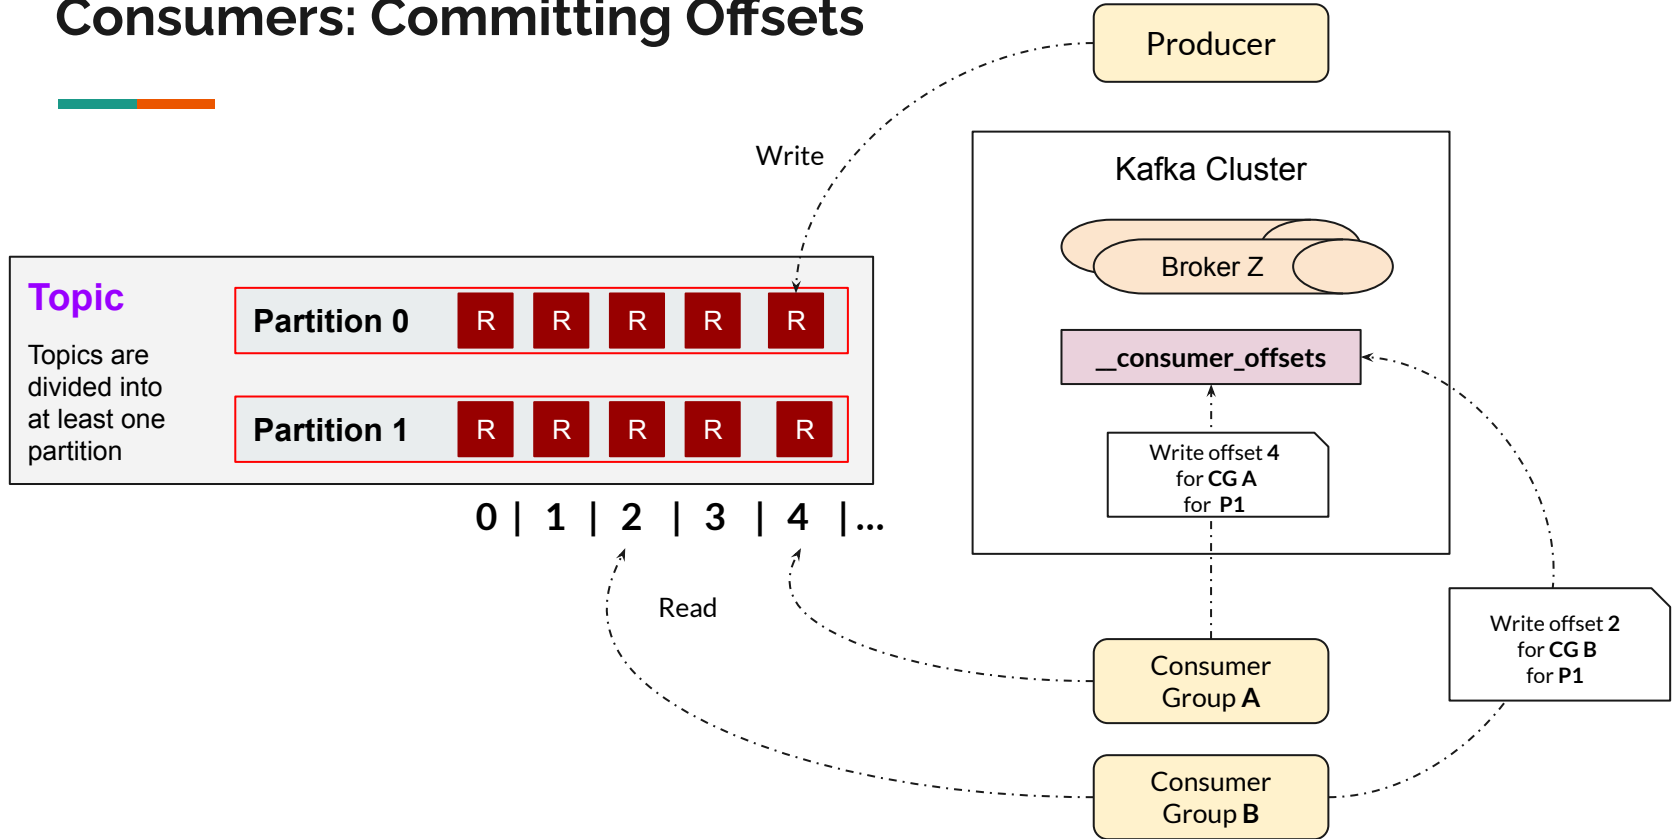
Brokers' Health

- Load Skewness (number of partitions on a broker)
*kafka.server:name=PartitionCount,
type=ReplicaManager*
- Log Flush Latency
*kafka.log:name=LogFlushRateAndTimeMs,
type=LogFlushStats*
- Network Request & Error Rate
- Fetcher Lag (per topic per partition)
*kafka.server:name=ConsumerLag,
type=FetcherLagMetrics,clientId=[-\\.w]+,topic=[-\\.w]+,
partition=[0-9]+)*

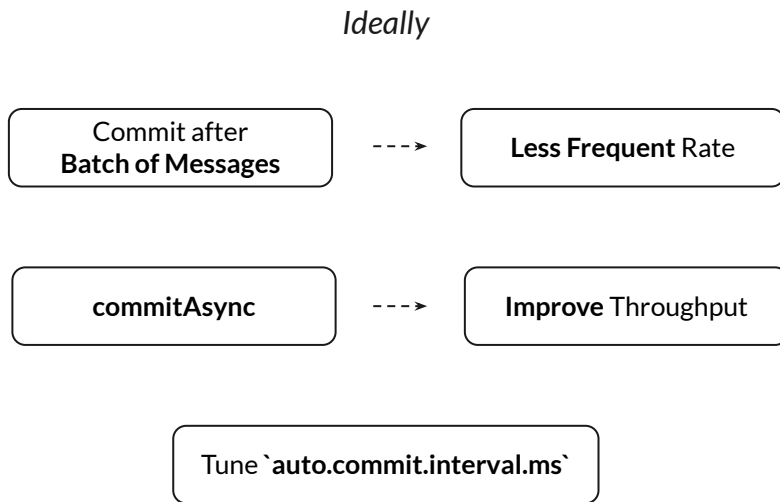
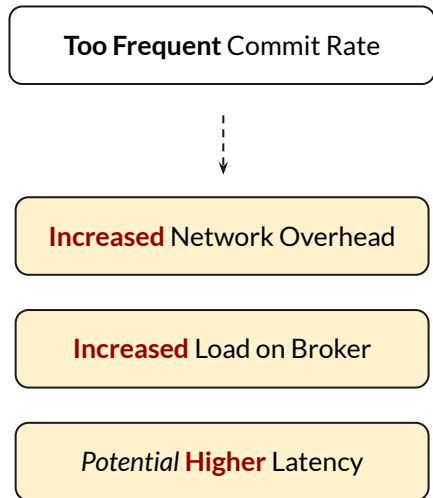
Brokers' Health



Consumers: Committing Offsets



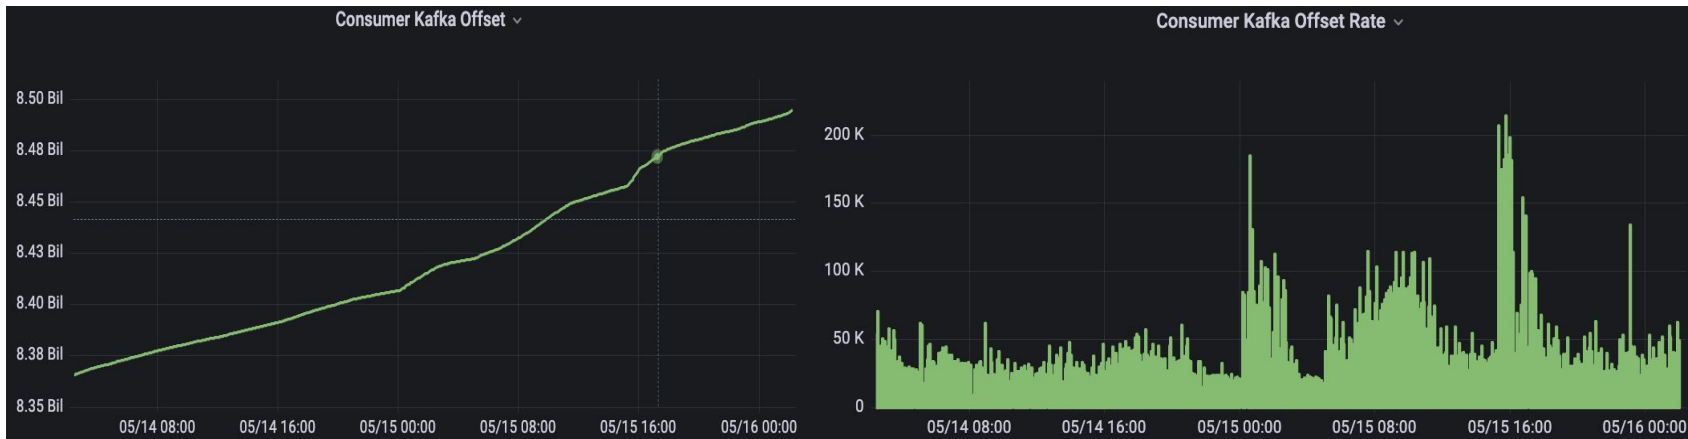
Consumers: Commit Rate



Consumption Rate

`__consumer_offsets` topic can be consumed;
offset long value emitted as metric to visualise **every consumer-partition's committed offset** in near real-time.

Reference [Burrow](#) (discussed later) which already does it.



Lag, Evaluations & Alerts: Burrow

Monitoring tool provides **consumer lag check as a service**

Exposes **offset lag** for all consumer-partition combination as **Prometheus metrics**.



Monitors committed offsets and **calculates the status of those consumers** on demand.

- Able to send alerts

Lag - Offsets Trend Evaluation

Lag = Head Offset of the Broker - Consumer's Offset

Store metrics locally in TSDB format. Evaluation runs on a sliding window periodically.



Status of consumer: OK

Lag Series with no Uptrend & Consumer Offset Series not Stalled


Lag - Consumer is Slow!



Status of consumer: **WARNING**

Lag Series with Uptrend & Consumer Offset Series not Stalled

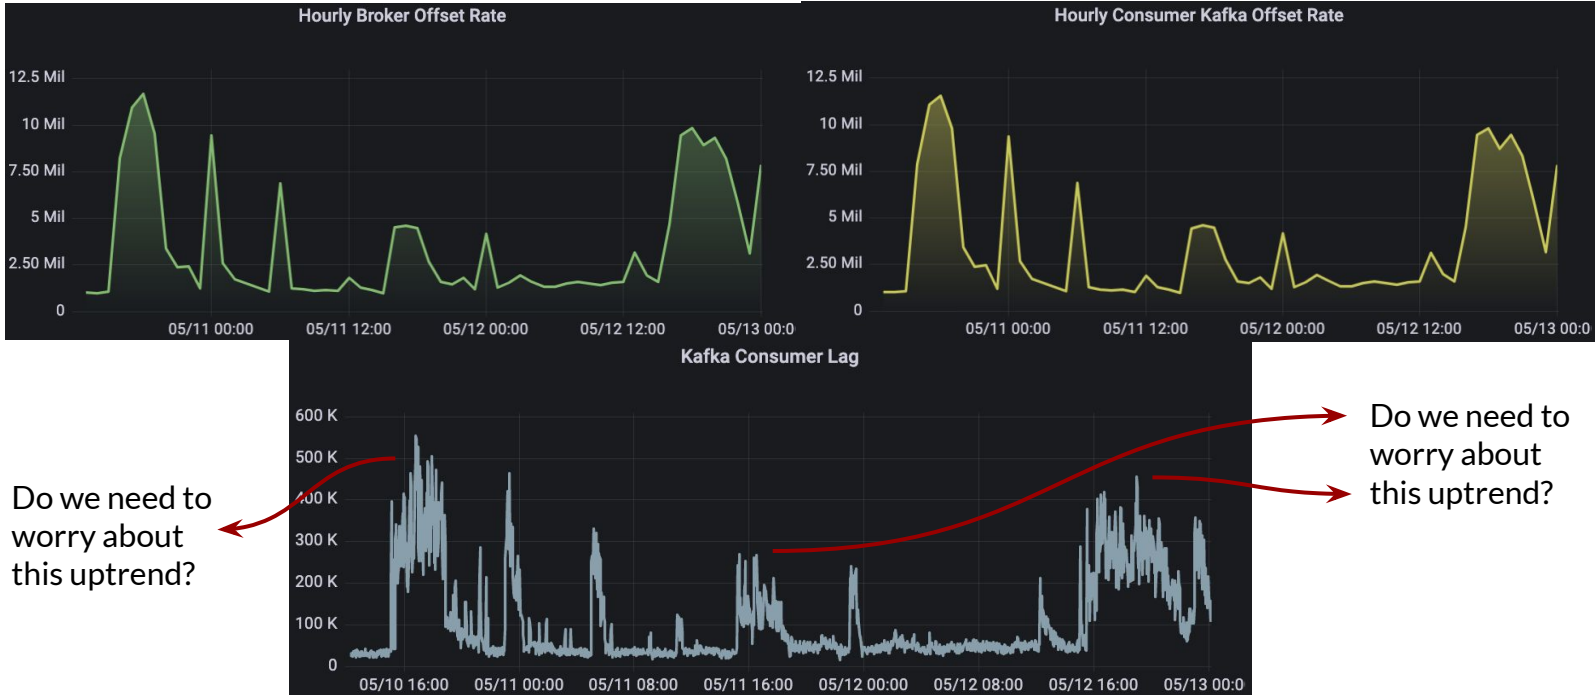
Lag - Consumer Stalled



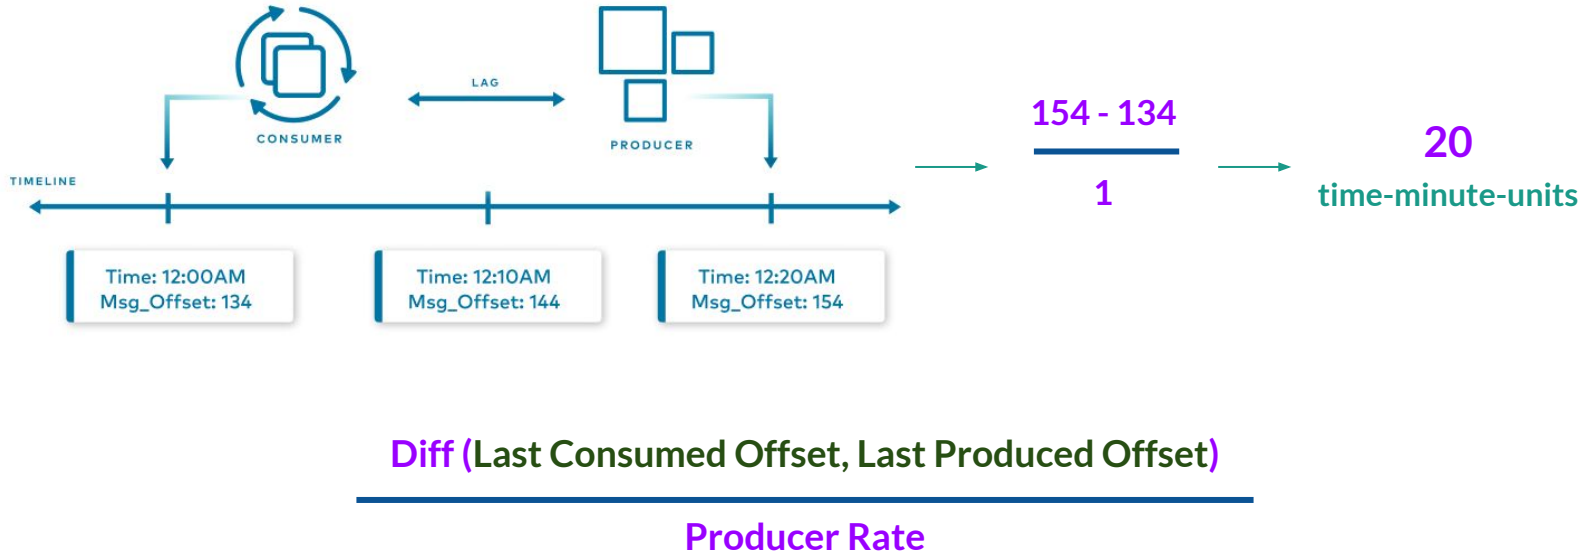
Status of consumer: **STALLED**

Lag Series with Uptrend & Consumer Offset Series Stalled

Lag - Observability




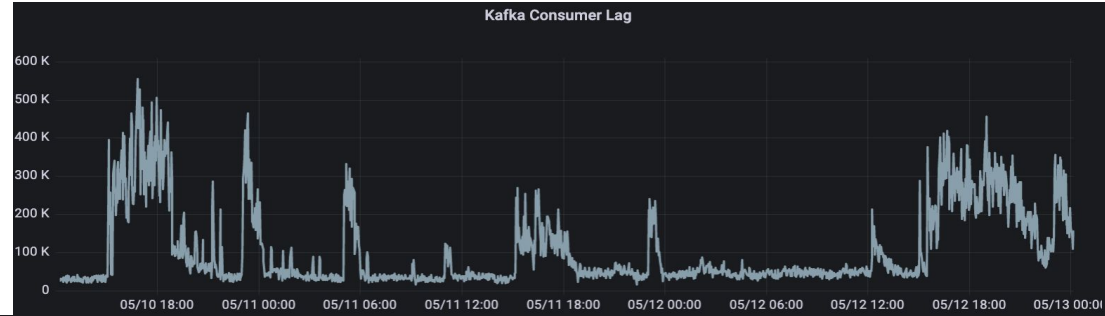
Lag - Time Based



Lag - Time Based



Normal Trend!



Trend Analysis



Keep track of high-level metrics for:

- Rate of Topic Growth → Do we need more partitions?
- Weekly / Monthly / Periodic
Producer / Consumer Rate → Keeping tabs on abnormal spikes!
- TTL / Retention data long enough to hold
data for consumption → If time lag for consumer-partition
goes beyond control!
- Infrastructure supporting Kafka cluster
requirements → CPU, Memory, Network, IO capacity,
GC Activity
- *Zookeeper supporting Kafka cluster
requirements* → *How many topics / partitions state can
be kept?*

* ZK has been deprecated / removed in newer versions

References



- Kafka Documentation - <https://docs.confluent.io/platform/current/kafka/monitoring.html>
- Kafka Architecture - <https://www.projectpro.io/article/apache-kafka-architecture-/442>
- Blog Posts & Talks
 - Peppredata - <https://www.youtube.com/watch?v=R6OKibnXpBs&t=385s>
 - <https://medium.com/quantyca/how-to-monitor-your-kafka-cluster-efficiently-d45ce37c02f1>
 - Stephan Meraak - <https://www.youtube.com/watch?v=XXLe0KNEbR4>
- Confluent Center - <https://docs.confluent.io/platform/current/control-center/index.html>
- Burrow - <https://github.com/linkedin/Burrow>
- Prometheus (<https://prometheus.io/>) & Grafana (<https://grafana.com/>)
- Need for Observability
 - <https://cloud.google.com/architecture/devops/devops-measurement-monitoring-and-observability>



Thank You!
@sarkaramrit2